

Health, Trigger & Retention

Automation That Surfaces Problems, Not Replaces Judgment

Purpose

Protect against mass churn in a segment where no single account is large enough to individually justify attention, but the aggregate is a real business risk if it silently erodes.

Health Score vs. Triggers — the Hierarchy

The health score identifies overall account condition. Triggers identify specific events that require action. A trigger can override the composite score — a single severe event (a critical ticket, a payment failure) escalates an account regardless of what the broader score says, since the score is a lagging composite and a trigger is a real-time signal.

Health Score Governance

Health-score weights should be reviewed quarterly against actual churn, contraction, and recovery outcomes, with documented human overrides where context materially contradicts the score. In practice, this means: - **Review cadence:** weights reviewed quarterly, not left static indefinitely - **Approval:** changes approved by CS leadership, not made unilaterally by whoever owns the platform config - **False positive/negative tracking:** each quarter, check which accounts the score got wrong — flagged as at-risk but stayed healthy, or scored healthy but churned — and adjust weighting accordingly - **CSM override:** a pod owner can override the automated score when direct account knowledge materially contradicts it, but the override and its reasoning must be logged, not silent

Automation — Three Layers

1. **Lifecycle automation** — welcome sequence, setup-stall nudges, feature-adoption emails at 30/60/90 days, renewal reminder sequence
2. **Health scoring** — rules-based score built from login frequency, feature adoption breadth, ticket volume/severity, and payment status; refreshed continuously
3. **Triggered outreach** — the score crossing a defined threshold auto-generates a task for a human; automation surfaces the problem, it never replaces the judgment call on what to do about it

Risk & Growth Triggers

Signal	Response
Usage drop / no login 14-30 days	Automated nudge, escalating to a human check-in past 30 days
Support ticket flagged urgent, or filed twice on the same issue	Immediate pod-owner review
Usage nearing plan/seat limits	Automated upsell prompt; human review only if engagement is also high
Multiple stakeholders logging in from one account	Flagged to Sales for qualification — possible signal the account is outgrowing Tier 4
Repeated use by a previously inactive stakeholder group, or a request for a higher-package feature	Flagged to Sales for qualification — possible expansion signal

Renewal Branch — the 90-Day Decision Point

Automated 90/60/30-day reminder sequence, but the health score at the 90-day mark decides the actual path: - **Healthy, auto-renew-eligible account:** automated renewal path continues, with confirmation and exception monitoring, no bespoke save motion required - **Degraded → routed to the pod owner** for a manual save attempt *before* the renewal date, not after

Skipping this branch is the most common design mistake at this tier — a flat, fixed-date reminder wastes a human touch on healthy accounts and silently loses struggling ones to an automated email they never engaged with.

Contraction & Cancellation at Volume

Same underlying logic as Tiers 1-3, run through the automated system rather than individual CSM judgment: a renewal quote below current value, or a non-renewal past the auto-renew window, fires the same trigger machinery. The pod owner runs any save attempt; anything requiring a discount routes to Sales. Exit data at this tier is captured via a lightweight 2-question exit survey rather than a live conversation — cheap enough to run at volume, and it's the only real source of pattern data on why a cohort is churning in aggregate.

Escalation Rule

Escalation is tiered by severity, not a flat count. One severe trigger prompts immediate review. Two moderate triggers prompt manual review. Repeated low-severity triggers prompt a pattern review. All inside a rolling 60-day window, regardless of the automated score — pattern-level risk overrides waiting for automation to resolve it on its own.

Measuring Stick

Metric	What it tells you
Activation rate	Is onboarding actually working
Time-to-first-value	Is the self-serve path fast enough
Segment churn rate (quarterly)	The core health number
Cost-to-serve per account	Is automation actually cheap, or just invisible
Trigger-response time (SLA)	Is the safety net working, not just existing
Escalation rate to human	Too low can mean problems are being missed, not resolved
Office hours / cohort attendance	Are the opt-in human touches actually being used
Knowledge base deflection rate	Is self-serve content working, or is the bot dead-ending people

Evergreen Contract Note

Evergreen removes the one built-in checkpoint a term contract provides. A mandatory quarterly health review, per account or per cohort, is the substitute checkpoint — built in specifically because no contract date is forcing one.

Artifacts

- Health scoring formula and weighting document
- Trigger rule set (embedded in the CS platform)
- 90-day renewal branch logic
- 2-question exit survey template
- Quarterly segment-level reporting template